



CSCI 3753: Operating Systems Fall 2024

Dylan Sain

Department of Computer Science

University of Colorado Boulder

Programming Assignment 1

Q & A



Week 3: Loadable Kernel Modules & Input/Output in C

How to Extend/Interact with Kernel?

- Method 1: System calls

Recompile the kernel !!!



- Method 2: Loadable Kernel Module (LKM)
 - An object file that contains a chunk of code to add to the base kernel of an operating system while it is **RUNNING**.
 - Typically, it is used to add support for *device drivers, filesystem drivers, etc.*

LKM Pros & Cons

- Pros:

- DON'T have to REBUILD the kernel
- Be MUCH faster to maintain and debug
- Save memory
 - Loaded only when we are actually using them
 - Unloaded in order to free memory and other resources when it is no longer required

- Cons:

- FRAGMENTATION penalty
- Security

LKM Command on Terminal

- **lsmod**: List currently loaded LKMs.
- **insmod**: Insert an LKM into the kernel.
- **rmmod**: Remove an LKM from the kernel.
- **modprobe**: Insert/remove an LKM or set of LKMs intelligently.
 - e.g., if you must load A before loading B, modprobe will automatically load A when you tell it to load B.
- **kerneld**: Kernel daemon program
 - allows kernel modules to be loaded automatically rather than manually with insmod/modprobe
- **depmod**: Determine interdependencies between LKMs

A large teal arrow pointing to the right, with a fine diagonal line pattern, serving as a background for the title.

Exercise 1: lsmod

Test out lsmod! Do you see your function from PA1?

LKM command – *insmod*

- *insmod* makes an `init_module` system call to load the LKM into kernel memory.
- `init_module` system call invokes the LKM's initialization routine (`module_init`) right after it loads the LKM.



LKM command – ***rmmod***

- ***rmmod*** makes a `delete_module` system call to unload the LKM into kernel memory.
- `delete_module` system call invokes the LKM's cleanup routine (`module_exit`) right before it unloads the LKM.

Input-Output in C



Different Ways to Access I/O for File in C

- I/O streams library (C standard library functions)
 - Associate with FILE pointer

BUFFER

A user-space array maintained by the C I/O streams library.

- I/O system calls
 - Associate with a file descriptor

I/O stream Library vs I/O System Calls

I/O Stream Library


- More high level
- Creates a pointer to the file stream
- Has a file structure and buffered I/O
- Built on top of I/O Systems Calls

I/O System Calls


- Low level interactions
- Provides a file descriptor
- Useful for working with other lower level system calls
- FDs have many different functions that go beyond the stream library

I/O for File Terminology

- **File descriptor (fd)**
 - Integer that uniquely identifies an open file of the process
- **File descriptor table**
 - The collection of integer array indices that are file descriptors in which elements are pointers to file table entries.
 - One unique file descriptors table is provided in operating system for each process.
- **File table entries**
 - A structure in-memory surrogate for an open file, which is created when process request to opens file and these entries maintains file position.



Standard File Descriptors



- When any process starts, that process file descriptors table's fd 0, 1, 2 open automatically.
- By default, each of these 3 fd references file table entry for a file named **`/dev/tty`**
 - **Read from stdin => read from fd 0**: Whenever we write any character from keyboard, it read from stdin through fd 0 and save to file named `/dev/tty`.
 - **Write to stdout => write to fd 1**: Whenever we see any output to the screen, it's from the file named `/dev/tty` and written to stdout in screen through fd 1.
 - **Write to stderr => write to fd 2**: We see any error to the screen, it is also from that file write to stderr in screen through fd 2.

Different Ways to Access I/O for File in C

	I/O Stream Library	I/O System Calls
Create a new file	<code>fopen()</code>	<code>creat()</code>
Open a file	<code>fopen()</code>	<code>open()</code>
Close a file	<code>fclose()</code>	<code>close()</code>
Read from a file	<code>fscanf()</code> , <code>getc()</code> , <code>getw()</code> , <code>fread()</code>	<code>read()</code>
Write to a file	<code>fprintf()</code> , <code>putc()</code> , <code>putw()</code> , <code>fwrite()</code>	<code>write()</code>
Sets the position of a file pointer to a specified location	<code>fseek()</code>	<code>lseek()</code>

fopen() vs open()

FILE* fopen(const char *filename, const char *mode)

- #include <stdio.h>
- Parameters:
 - filename : path to file
 - mode : "w", "r", "a", "w+", "r+", "a+"

int open(const char* Path, int flags[, mode_t mode])

- #include <fcntl.h>
- Parameters:
 - path : path to file
 - flags : O_RDONLY, O_WRONLY, O_RDWR, O_CREAT, etc.
 - Mode : indicates permissions of new file

Reading & Writing in I/O Stream

I/O Stream Function	Description
<code>fprintf()</code> <code>fwrite()</code> <code>fputs()</code>	Write a block of data to a file
<code>fscanf()</code> <code>fread()</code> <code>fgets()</code>	Read a block of data from a file
<code>fgetc()</code>	Read a single character from a file
<code>fputc()</code>	Write a single character to a file
<code>ftell()</code>	Returns the current position of a file pointer
<code>rewind()</code>	Sets the file pointer at the beginning of a file

Reading & Writing in I/O Standard Stream

I/O Stream Function	Description
<code>getc()</code>	Reads a single character from stdin
<code>putc()</code>	Writes a single character to stdout
<code>getw()</code>	Reads an integer from stdin
<code>putw()</code>	Writing an integer to stdout
<code>scanf()</code> <code>getline()</code> <code>gets()</code>	Reads a block of data from stdin
<code>printf()</code>	Writing a block of data to stdout

Appendix

I/O System calls

- **creat()**: create a new empty file
- **open()**: open the file for reading, writing or both
- **close()**: close the file which pointed by fd
- **read()**: read n bytes of input from the file indicated by the file descriptor fd into the memory area indicated by buf
- **write()**: write n bytes from buf to the file or socket associated with fd
- **lseek()**: change the location of the read/write pointer of the fd

I/O System calls

- **creat()**: create a new empty file

- Syntax

- ```
int creat(char *filename, mode_t mode)
```

- Parameters:

- filename : name of the file which you want to create
    - mode : indicates permissions of new file

- How it works in OS

- Create new empty file on disk
  - Create file table entry
  - Set first unused file descriptor to point to file table entry
  - Return file descriptor used or -1 upon failure

# I/O System calls

- **open()**: open the file for reading, writing or both

- Syntax

```
int open(const char* Path, int flags[,
 mode_t mode])
```

- Parameters:

- path : path to file
- flags : O\_RDONLY, O\_WRONLY, O\_RDWR, O\_CREAT, etc.
- mode: indicates permissions of new file

- How it works in OS

- Find existing file on disk
- Create file table entry
- Set first unused file descriptor to point to file table entry
- Return file descriptor used or -1 upon failure

# I/O System calls

- **close()**: close the file which pointed by fd

- Syntax

```
int close(int fd)
```

- Parameters:

- fd: file descriptor

- How it works in OS

- Destroy file table entry referenced by element fd of file descriptor table as long as no other process is pointing to it !!!
- Set element fd of file descriptor table to NULL
- Return 0 on success or -1 on error

# I/O System calls

- **read()**: read bytes of input from fd into memory area
  - Syntax

```
size_t read(int fd, void* buf, size_t n)
```
  - Parameters:
    - fd: file descriptor
    - buf: buffer to read data from
    - cnt: length of buffer
  - Returns:
    - 0 on reaching end of file
    - -1 on error
    - -1 on signal interrupt
    - Number of bytes read on success



# I/O System calls

- **write()**: write bytes of input from memory area to fd
  - Syntax

```
size_t write(int fd, void* buf, size_t n)
```
  - Parameters:
    - fd: file descriptor
    - buf: buffer to read data from
    - cnt: length of buffer
  - Returns:
    - 0 on reaching end of file
    - -1 on error
    - -1 on signal interrupt
    - Number of bytes written on success

# I/O System calls

- **lseek()**: change the location of the read/write pointer of the fd
  - Syntax

```
off_t lseek(int fildes, off_t offset, int whence)
```
  - Parameters:
    - fildes: fd of the pointer that is going to be moved
    - offset: offset of the pointer (measured in bytes)
    - whence : The method in which the offset is to be interpreted
  - Returns:
    - the offset of the pointer (in bytes) from the beginning of the file.
    - -1 on an error moving the pointer