



# CSCI 3753: Operating Systems Fall 2024

Dylan Sain

Department of Computer Science

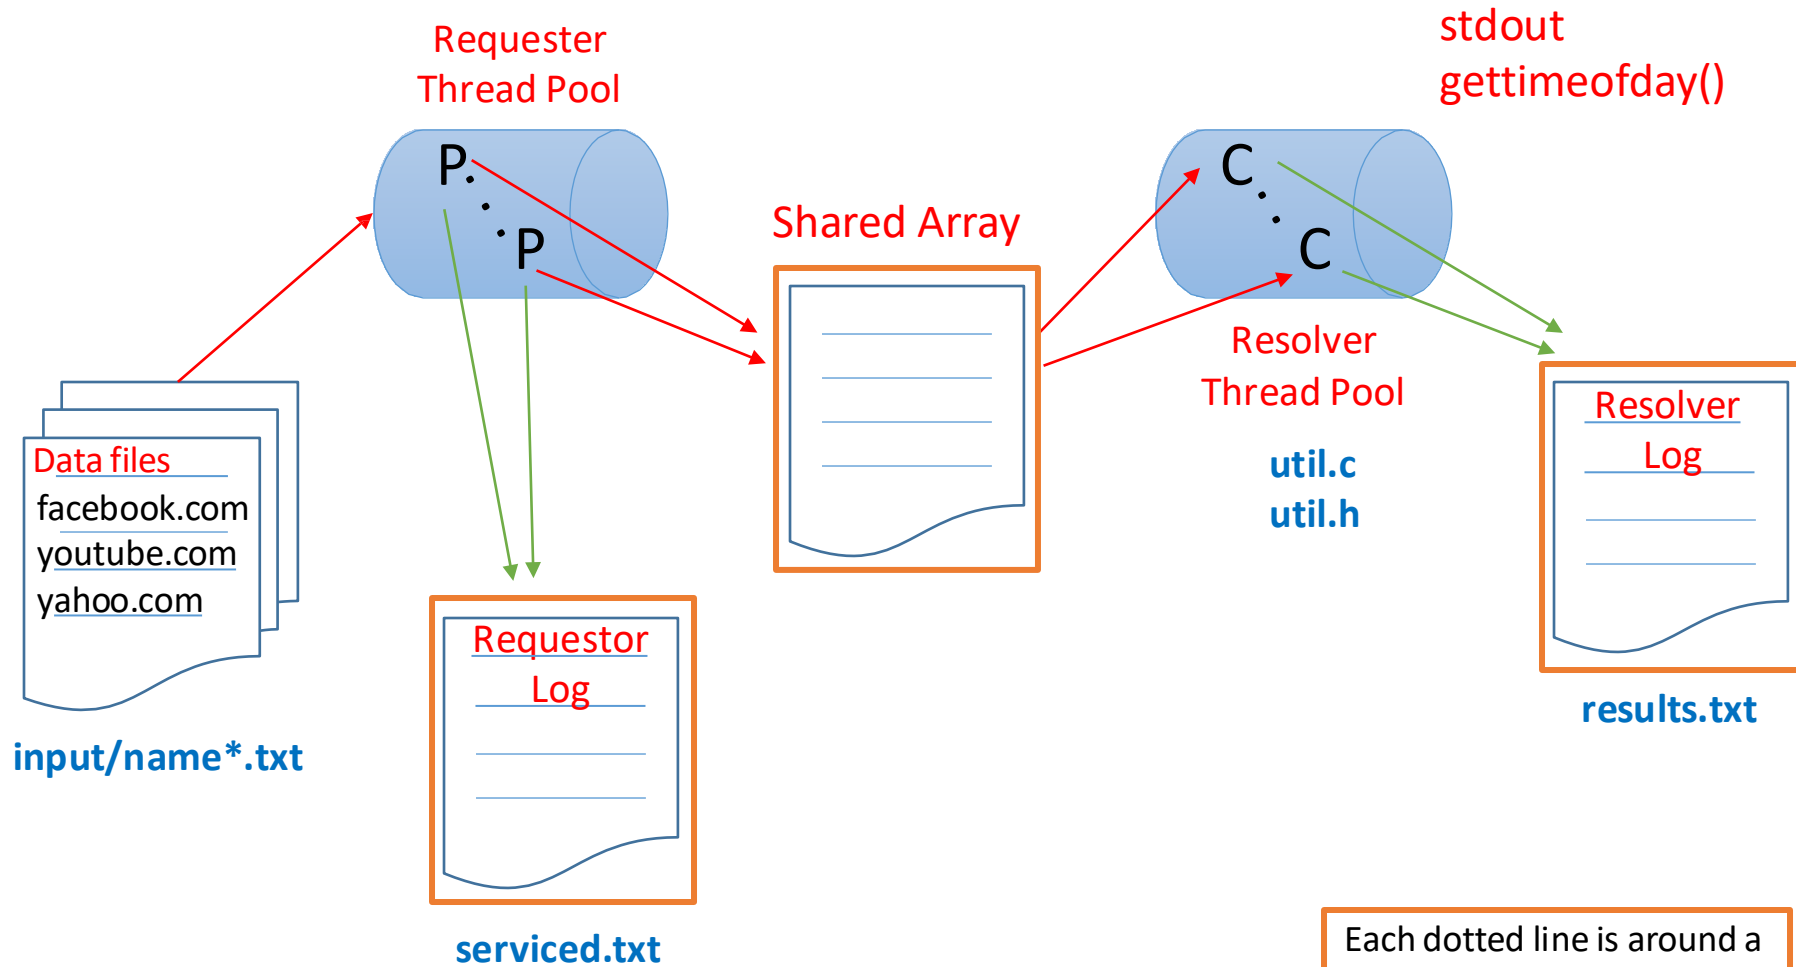
University of Colorado Boulder

# Week 10: PA6 and Process Scheduling



# Overview PA6


Total runtime is ...



Each dotted line is around a shared resource that must be protected from race conditions.

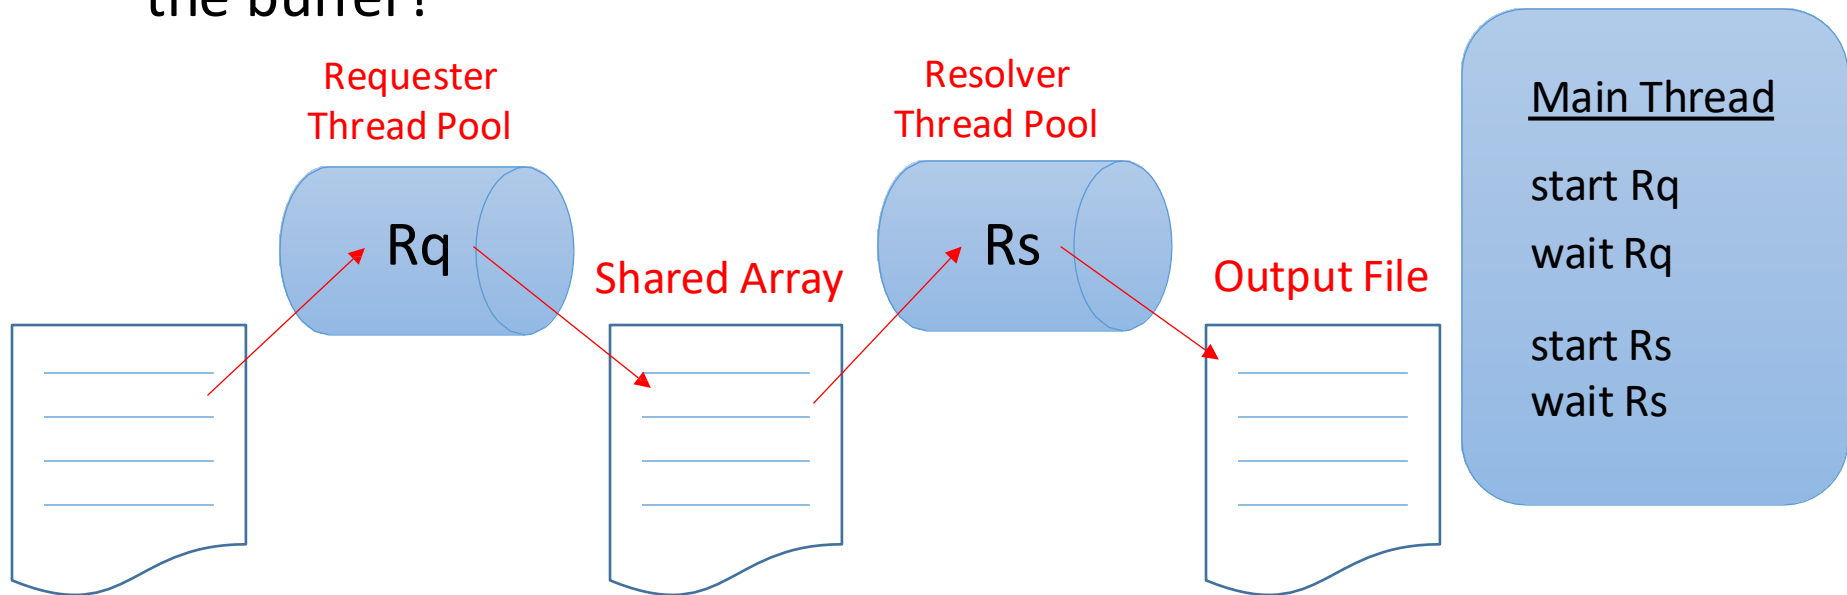
# Implementation – Step 1

- **Action:** Create a simple program to create a requester thread that will repeatedly read a line from a given file and add an entry into the shared buffer
- **Validation:** Does the buffer have the correct number of entries?



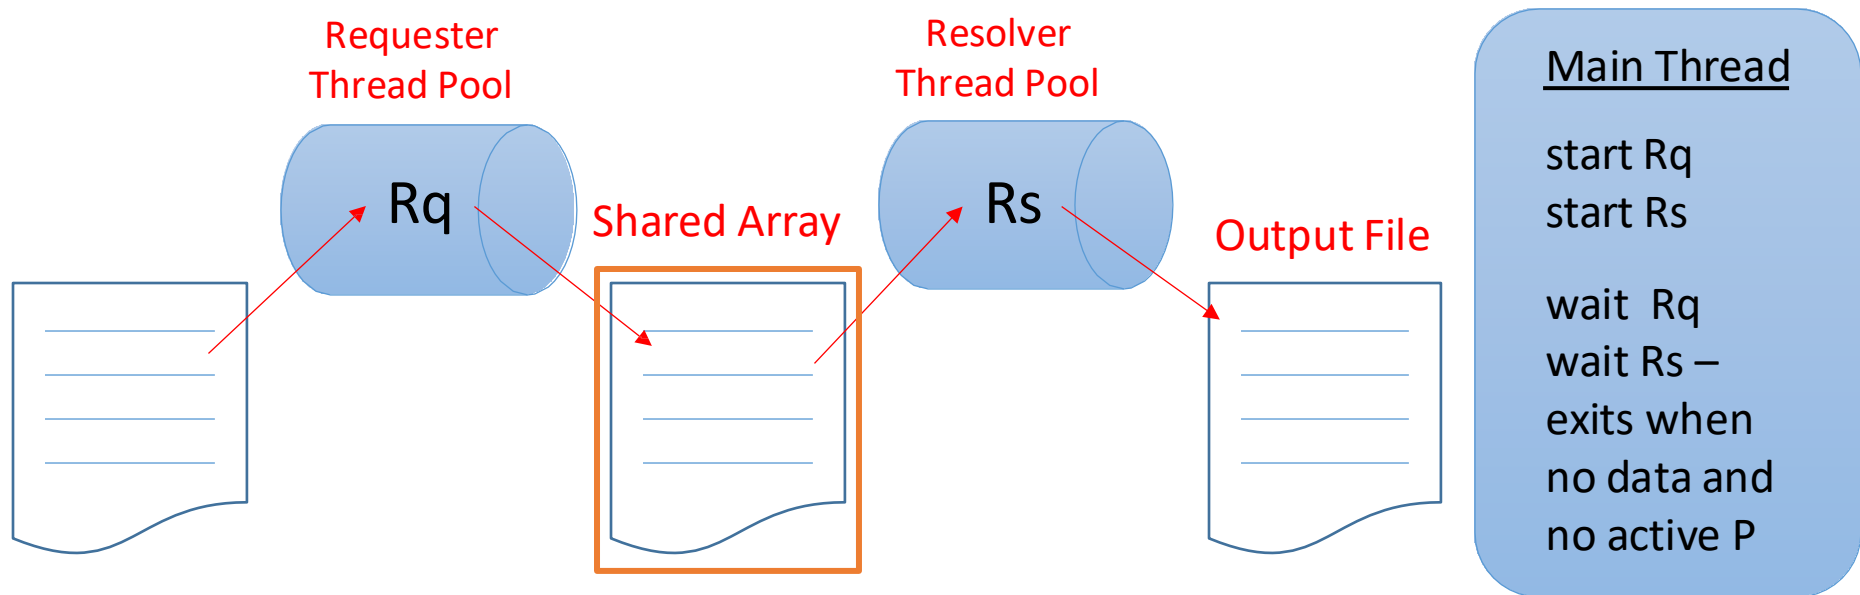
# Implementation – Step 2

- **Action:** Use the result in step 1, once that process is complete, start a resolver thread to take items out of the buffer. Then, write the results to an output file.
- **Validation:** Does the output file contain the entries stored in the buffer?



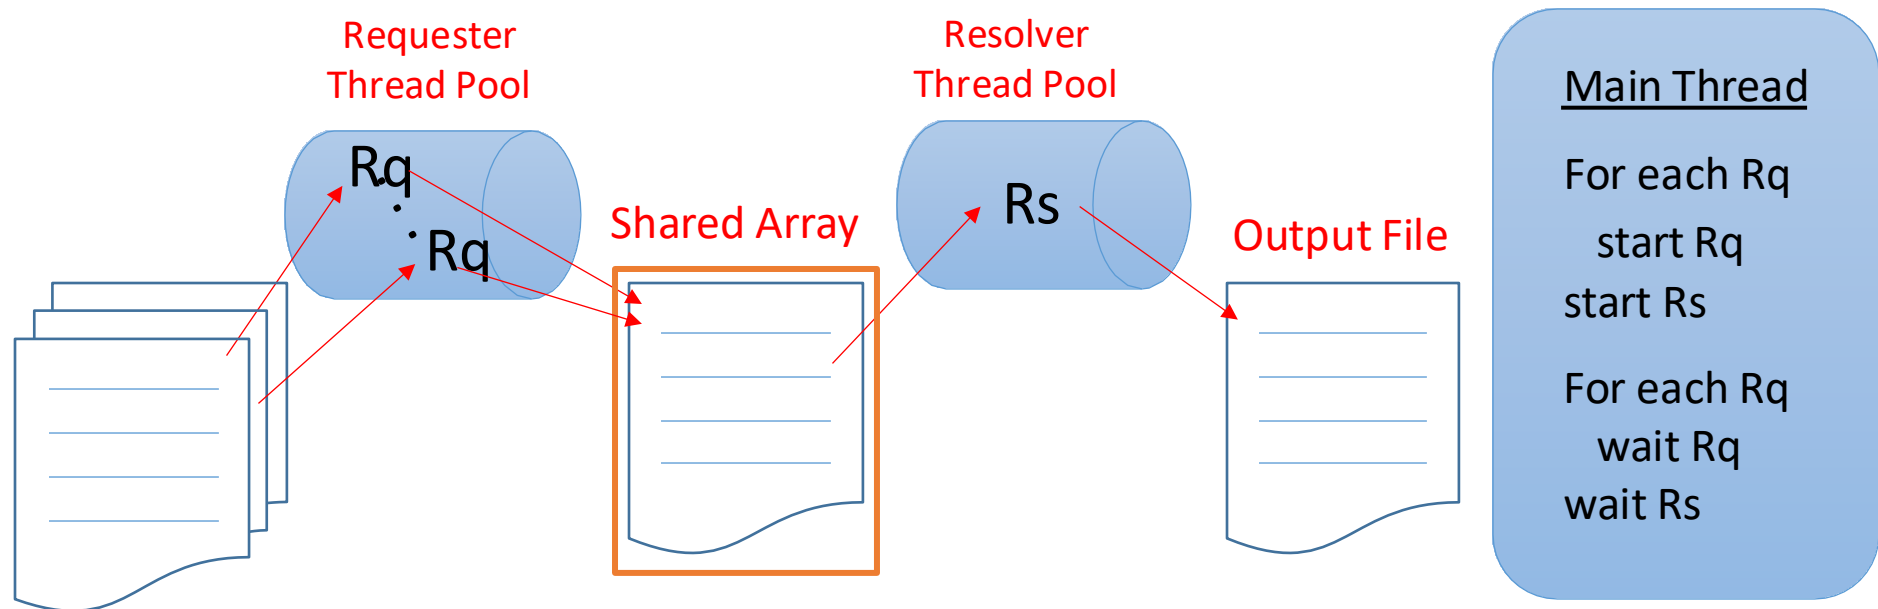
# Implementation – Step 3

- **Action:** Once you are assured that your application can write and read to the buffer correctly (although serially), then try to make them run concurrently. Multiple processes accessing and modifying the same data can cause race conditions. You must protect the critical sections of each thread with a mutex.



# Implementation – Step 4

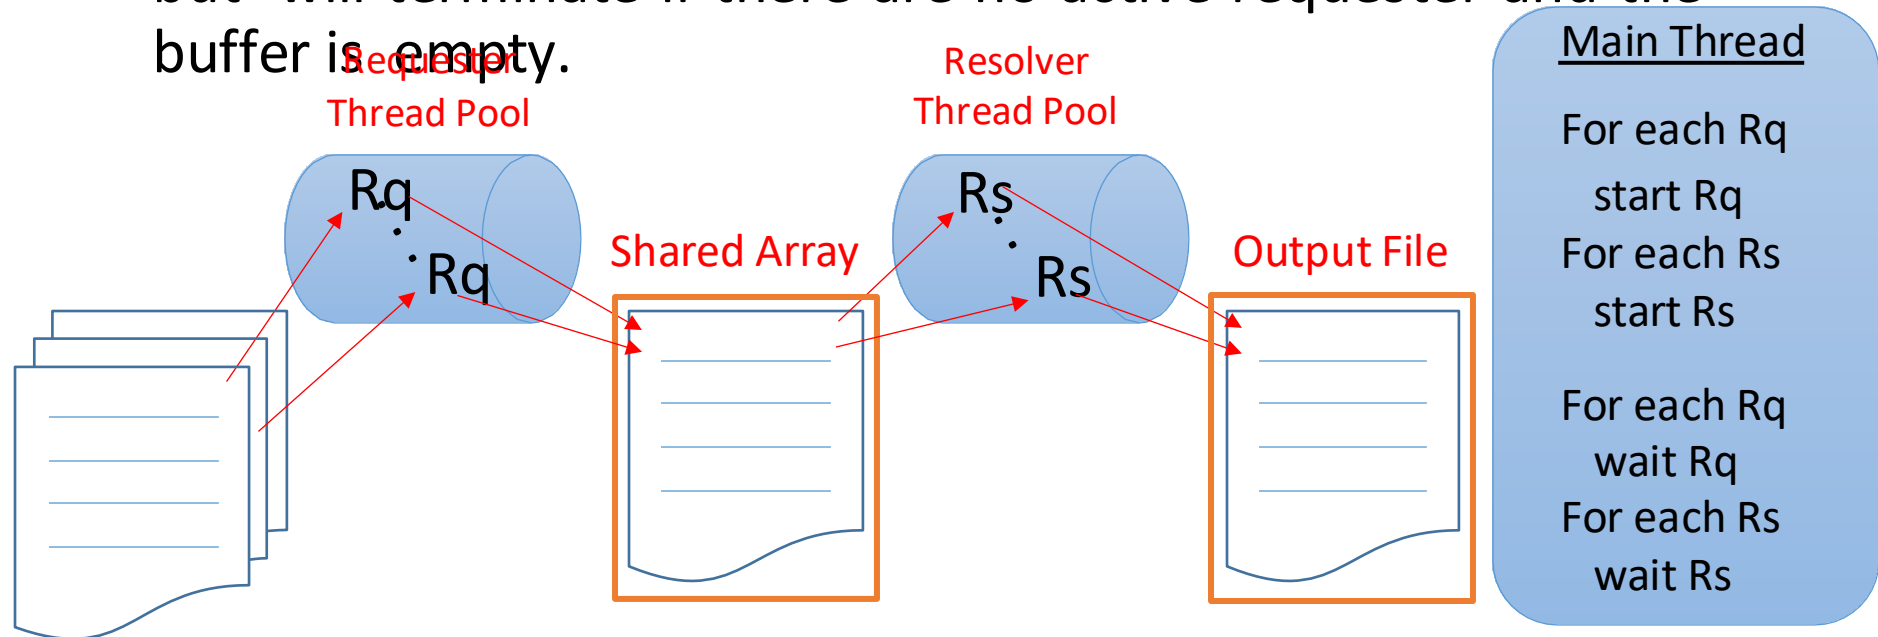
- **Action:** create multiple requester threads to read from multiple different files. Each requester can read single lines from a different file. The requester will terminate when all lines from the file have been processed.



# Implementation – Step 5

- **Action:** create multiple resolver threads to read from multiple requester threads via a single shared buffer.

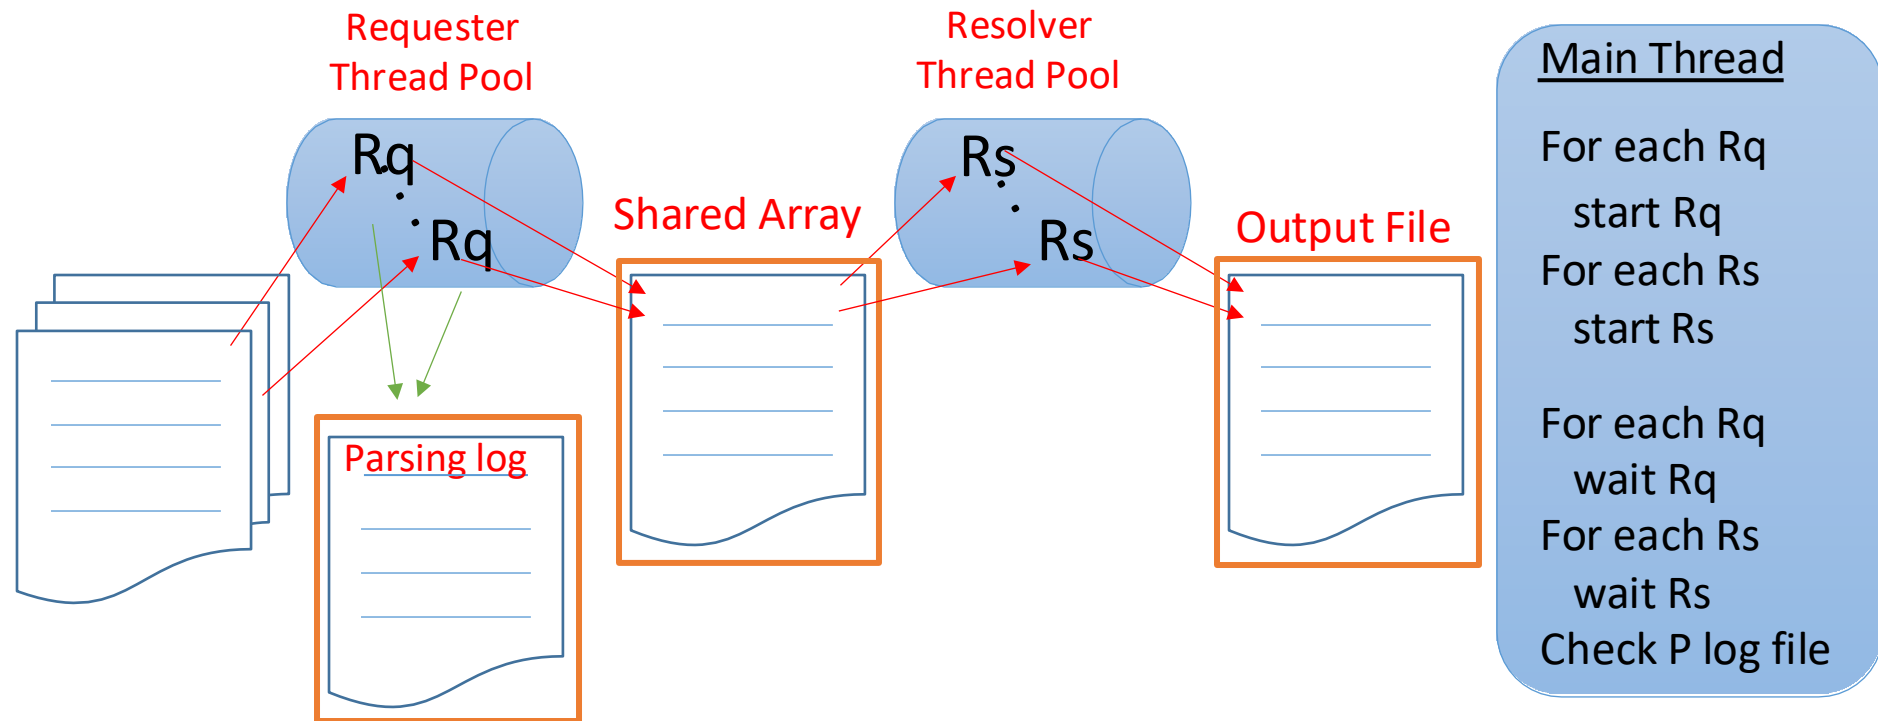
The resolver will wait for data (spin wait is acceptable) but will terminate if there are no active requester and the buffer is empty.





# Implementation – Step 6

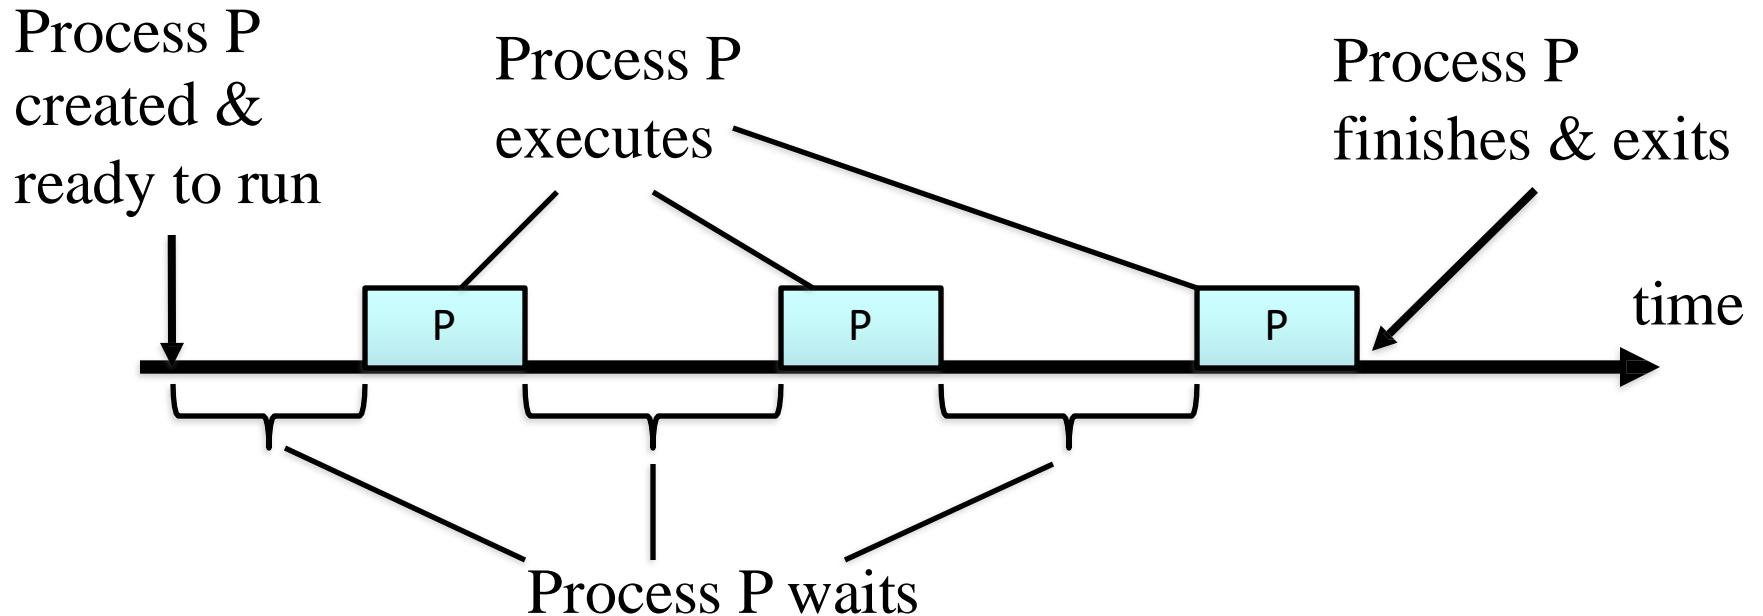
- **Action:** move back to the requester threads, each thread must record the data it has processed. Files are a shared resource and therefore must be protected from multiple processes accessing it.



# Scheduling



# Per Job or Task Metrics



- Execution time
- Wait time
- Turnaround time
- Response Time

A large teal arrow pointing to the right, with a fine diagonal line pattern, serving as a background for the title text.

# Exercise: Test your skills!

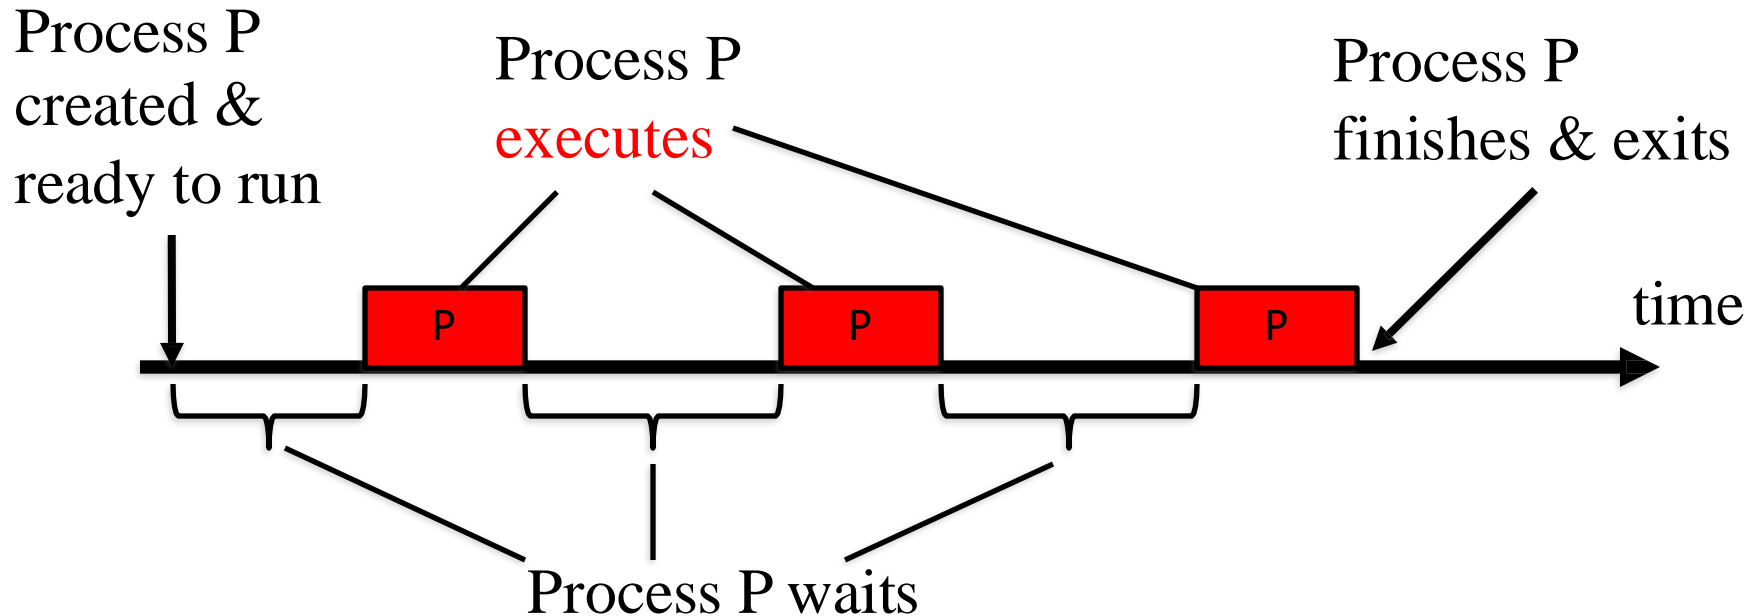
Go to Kahoot.it

Use the code

Add your name (please  
use your real name)


Choose the correct  
answer!

# Per Job or Task Metrics




- **Execution time**  $E(P)$  = the time on the CPU required to **fully execute** process P
  - Sum up the time slices given to process P
  - Also called the “**burst time**” by textbook

# Per Job or Task Metrics



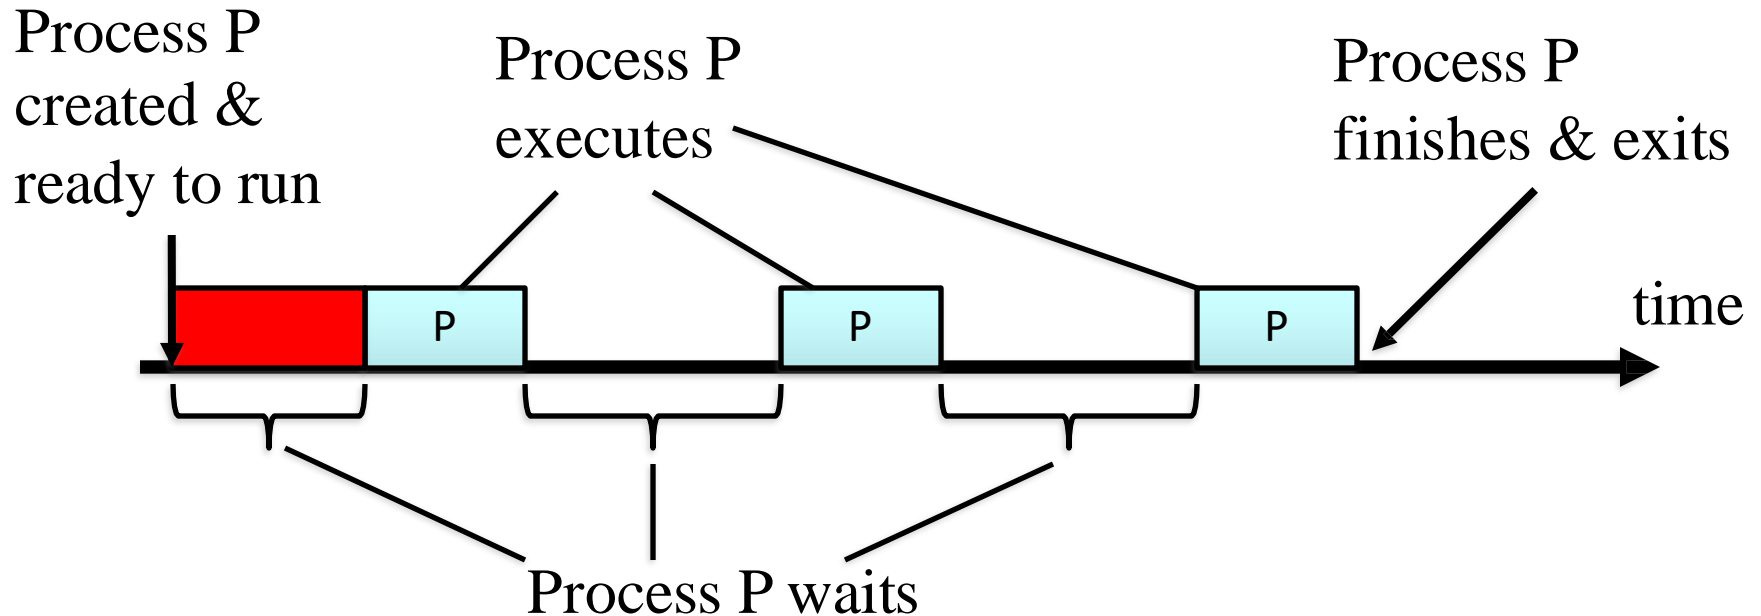
- **Wait time**  $W(P)$  = the time process P is in the ready state/queue waiting but not running
  - Sum up the gaps between time slices given to process P
  - But, does **NOT** include I/O waiting time

# Per Job or Task Metrics



- **Turnaround time**  $T(P)$  the time from 1<sup>st</sup> entry of process P into the ready queue to its final exit from the system (exits last run state)
  - Does include time waiting and time for IO to complete

# Per Job or Task Metrics



- **Response time**  $R(P)$  = the time from 1<sup>st</sup> entry of process P into the ready queue to its 1<sup>st</sup> scheduling on the CPU (1<sup>st</sup> occurrence in running state)
  - Useful for interactive tasks



# Scheduling Policies

- First Come First Serve (FCFS) Scheduling
- Shortest Job First (SJF) Scheduling
- Round Robin Scheduling
  - Weighted RR
- Earliest Deadline First Scheduling
  - W/ or w/o preemption
  - Least slack (slack time = time until deadline – remaining execution time)
- Priority-based or Multilevel Queue Scheduling
- Multilevel Feedback Queue Scheduling
- Completely Fair Scheduler in Linux

# Scheduling Policies


- **First Come First Serve (FCFS)** Scheduling

- Average wait time?

→  $(6+14+24)/4 = 11 \text{ ms}$

- Average turnaround time?

→  $(6+14+24+27)/4 = 17.75 \text{ ms}$




| Task | CPU Execution Time (ms) |
|------|-------------------------|
| T1   | 6                       |
| T2   | 8                       |
| T3   | 10                      |
| T4   | 3                       |

- All processes arrived just before time 0 in the order.

# Scheduling Policies

- **Shortest Job First (SJF)** Scheduling

- Average wait time?  
→  $(3+9+17)/4 = 7.25$  ms
- Average turnaround time?  
→  $(3+9+17+27)/4 = 14$  ms



| Task | CPU Execution Time (ms) |
|------|-------------------------|
| T1   | 6                       |
| T2   | 8                       |
| T3   | 10                      |
| T4   | 3                       |

- All processes arrived just before time 0.

# Scheduling Policies


- **Shortest Job First (SJF) Scheduling w/ Preemption**

- Average wait time?

→  $(6+3+17)/4 = 6.5$  ms

- Average turnaround time?

→  $(6+6+3+8+3+17+10)/4 = 13.25$  ms



| Task | CPU Execution Time (ms) |
|------|-------------------------|
| T1   | 6                       |
| T2   | 8                       |
| T3   | 10                      |
| T4   | 3                       |

- T1, T2, and T3 arrived just before time 0. T4 arrived at 10ms.

# Scheduling Policies


- **Round Robin** Scheduling

- Let **time slice = 4 ms**
- Assuming a **1ms switching time**

- Average wait time?  
→  $(15+5+13+10+13+15)/4 = 17.75$  ms
- Average response time?  
→  $(5+10+15)/4 = 7.5$  ms

All processes arrived just before time 0 in the order.

| Task | CPU Execution Time (ms) |
|------|-------------------------|
| T1   | 6                       |
| T2   | 8                       |
| T3   | 10                      |
| T4   | 3                       |



# Scheduling Policies

- **Earliest Deadline First (EDF) w/ Preemption**
  - Let **time slice = 20 ticks**
  - Assuming a **5 ticks of switching time**

| Task | CPU Execution Time | Arrival Time | Deadline |
|------|--------------------|--------------|----------|
| T1   | 30                 | 0            | 100      |
| T2   | 90                 | 20           | 230      |
| T3   | 40                 | 55           | 145      |
| T4   | 20                 | 85           | 145      |

